

## Assignment 3, 2017

Released: 25 April.

Submit test data (team effort): Friday 5 May at 23:00

Submit compiler (team effort): Friday 19 May at 23:00

## Objectives

To build a better understanding of a compiler's back-end, code generation, symbol tables, run-time structures, and optimization. To practice cooperative, staged software development.

## Background and context

The task is to write a compiler for a procedural language, Snick. The compiler translates source programs to the assembly language of a target machine Brill. These programs can then be run on a provided Brill emulator.

In an earlier stage, you wrote a parser for Snick. You may choose to start from that parser, or alternatively, start from one that has been (or will be) made available. In either case, correctness of the compiler, including the parser, is your responsibility.

This final stage also involves the completion of semantic analysis, code generation, and (optionally) optimization. The implementation language must be in OCaml.

## The source language: Snick

Snick is already known to you from Stage 1, except that we refine the definition of well-formed Snick programs to say that, in an array component expression  $id[e_1, \dots, e_n]$ ,  $n$  must be greater than 0, that is,  $id[ ]$  is not allowed.

A Snick program lives in a single file and consists of a number of procedure definitions. There are no global variables. One (parameter-less) procedure must be named “main”. Procedure parameters can be passed by value or by reference.

The language has three types, namely *int*, *float*, and *bool*. (The first two are *numeric* types and allow for arithmetic and comparison operators. We do not consider the Boolean values to be ordered, but Boolean values can still be compared for equality = and !=.) It also has static arrays. The “write” command can print integers, floating point numbers, booleans and, additionally, strings.

## Syntax

The following are reserved words: **and**, **bool**, **do**, **else**, **end**, **false**, **fi**, **float**, **if**, **int**, **not**, **od**, **or**, **proc**, **read**, **ref**, **then**, **true**, **val**, **while**, **write**. The lexical rules were given in the specification for Stage 1.

The arithmetic binary operators associate to the left, and unary operators have higher precedence (hence for example, ‘- 5 + 6’ and ‘4 - 2 - 1’ both evaluate to 1). A float constant is a

sequence of one or more digits, followed by a decimal point, and another sequence of one or more digits. A Boolean constant is **false** or **true**. A string constant (as can be used by “**write**”) is a sequence of characters between double quotes. The sequence itself cannot contain double quotes or escape characters, except it may include “\n”, representing a newline character. A backslash not preceding an ‘n’ is just considered part of the string.

A Snick program consists of one or more procedure definitions. Each definition consists of (in this order):

1. The keyword **proc**.
2. A procedure header.
3. A procedure body.
4. The keyword **end**.

The header has two components (in this order):

1. An identifier—the procedure’s name.
2. A comma-separated list of zero or more formal parameters within a pair of parentheses (so the parentheses are always present).

Each formal parameter has three components:

1. A parameter passing indicator (**val** or **ref**).
2. A type (**bool**, **float** or **int**).
3. An identifier.

The procedure body consists of zero or more local variable declarations, followed by a non-empty sequence of statements. A variable declaration consists of one of the keywords **bool**, **float** or **int**, followed by an identifier, terminated with a semicolon. Before the semicolon may be a non-empty comma-separated list of integer intervals, the list enclosed in square brackets, to indicate that the identifier names an array. An integer interval is of the form **m..n**, where **m** and **n** are integer constants. The number of intervals corresponds to the array’s dimension. For example, **float point[0..4,2..3,0..1]** declares a three-dimensional array **point** consisting of 20 floats. There may be any number of variable declarations, given in any order.

Atomic statements have one of the following forms:

```
<id> := <expr> ;
<id> [ <expr-list> ] := <expr> ;
read <id> ;
read <id> [ <expr-list> ] ;
write <expr> ;
<id> ( <opt-expr-list> ) ;
```

where **<expr-list>** is a non-empty comma-separated list of expressions and **<opt-expr-list>** is a (possibly empty) comma-separated list of expressions. Composite statements have one of the following forms:

```

if <expr> then <stmt-list> fi
if <expr> then <stmt-list> else <stmt-list> fi
while <expr> do <stmt-list> od

```

where `<stmt-list>` is a non-empty sequence of statements, atomic or composite. (Note that semicolons are used to *terminate* atomic statements, so that a sequence of statements—atomic or composite—does not require any punctuation to *separate* the components.)

Expressions have one of the following forms:

```

<id>
<id> [ <expr-list> ]
<const>
( <expr> )
<expr> <binop> <expr>
<unop> <expr>

```

The list of operators is

```

or
and
not
= != < <= > >=
+ -
* /
-

```

`not` is a unary prefix operator, and the bottom “`-`” is a unary prefix operator (unary minus). All the other operators are binary and infix. All the operators on the same line have the same precedence, and the ones on later lines have higher precedence; the highest precedence being given to unary minus. The six relational operators are not associative. The seven remaining binary operators are left-associative. The relational operators yield Boolean values `true` or `false`, according as the relation is true or not.

The language supports comments, which start at a `#` character and continue to the end of the line. White space is not significant.

## Static semantics

We define a *scalar* to be an object that contains a single value. This value can be inspected and updated. Variables and array components  $id[e_1, \dots, e_n]$  are scalars. Snick does not allow an array *as a whole* to be assigned or passed as a parameter. Only *selective updating* is possible. In other words, if  $id$  has been declared as an array then  $id$  is illegal as an expression on its own. It follows that every valid expression has type `bool`, `int`, or `float`.

All defined procedures must have distinct names. A defined procedure does not have to be called anywhere, and the definition of a procedure does not have to precede the (textually) first call to the procedure. For each procedure, the number of actual parameters in a call must be equal to the number of formal parameters in the procedure’s definition.

A formal procedure parameter is treated as a local variable. The scope of a declared variable (or of a formal procedure parameter) is the enclosing procedure definition. A variable must be declared (exactly once) before used. Similarly a list of formal parameters must all be distinct. However, the same variable/parameter name can be used in different procedures. An expression

in an actual argument position where the parameter passing method is “by reference” must be a scalar (variable or array component).

For each integer interval  $m..n$  in an array declaration,  $m \leq n$  must hold. Array identifiers obey the same scope rules as variables. The number  $n$  of indices in an array component expression  $id[e_1, \dots, e_n]$  must agree with the dimension given by the array’s declaration. An array cannot be passed as a whole entity to a procedure, but an array component, such as  $a[3, 2]$  can be used as a formal parameter, passed by value or reference. Since the array bounds are known at declaration-time, semantic analysis *could* do limited bounds checking at compile-time, but there is no requirement to do this.

The language is statically typed, that is, each variable and parameter has a fixed type, chosen by the programmer. The type rules for expressions are as follows:

- The type of a Boolean constant is **bool**.
- The type of an integer constant is **int**.
- The type of a float constant is **float**.
- The type of an expression  $id$  is the variable  $id$ ’s declared type ( $id$  cannot denote an array).
- In an array component  $id[e_1, \dots, e_n]$ , each  $e_j$  must have type **int**,  $n$  must equal the number of dimensions given in the declaration of the array  $id$ , and the type of the whole expression is the type given in the declaration of  $id$ .
- Arguments of the logical operators, and their results, must be of type **bool**.
- The two operands of  $=$  must have the same type; the same goes for  $!=$ . The result is Boolean.
- The two operands of any other comparison operator must either have the same numeric type, or one must be **int** and the other **float** (in which case the compiler must then convert the integer to a float before the comparison). The result is Boolean.
- The two operands of a binary arithmetic operator must either have the same numeric type, or one must be **int** and the other **float** (in which case the compiler must convert the integer to a float before the comparison). The result type is **float**, unless both operands have type **int**, in which case the result type is **int**.
- The operand of unary minus is of type **int** or **float**, and the result type is the same as the operand’s type.

The type rules for statements are as follows:

- In assignment statements, the variable or array component on the left-hand side and the expression on the right hand side must have the same type, with one exception: an **int** expression may be assigned to a **float** variable (the compiler must convert the value to be assigned from **int** to **float**). Note that arrays can only be updated selectively, that is, only array *components* (and scalar variables) can be assigned to.
- Conditions in **if** and **while** statements must be of type **bool**.

- In procedure calls, the type of an actual parameter must be the type of the corresponding formal parameter, or (for passing “by value”), the actual parameter may be of type `int` and the formal parameter of type `float` (in which case the compiler must convert the integer to a float).
- `read` takes the name of a scalar (variable or array component), and `write` takes a well-typed expression *or* a string literal.

Every program must contain a parameter-less procedure named “main”.

## Dynamic semantics

Numerical variables are automatically initialised to 0, and Boolean variables to `false`. This extends to arrays. The evaluation of an expression  $e_1/e_2$  results in a runtime error if  $e_2$  evaluates to 0.

The logical operators are *strict*, that is, they evaluate all their operands fully, rather than using short-circuit evaluation. For example, `5 < 8 or 5 > 8/0` yields a runtime error rather than `true`.

If the program reads or writes outside the bounds of an array, the behaviour is undefined. (There is no requirement that out-of-bounds array indices are detected at runtime, though this may be implemented as an optional extension, as discussed below).

`write` prints `int`, `float` and `bool` expressions to `stdout` in their standard syntactic forms, with no additional whitespace or newlines. If `write` is given a string literal, it prints out the characters of the string to `stdout`, with `\n` resulting in a newline character being printed. Similarly, `read` reads `int`, `float` and `bool` literals from `stdin` and assigns them to variables or array components. If the user input is not valid, execution terminates.<sup>1</sup>

The procedure “main” is the entry point, that is, execution of a program comes down to execution of a call to “main”.

The language allows for two ways of passing parameters. Call by value (`val`) is a copying mechanism. For each parameter  $e$  passed by value, the called procedure considers the corresponding formal parameter  $v$  a local variable and initialises this local variable to  $e$ ’s value.

Call by reference (`ref`) does not involve copying. Instead the called procedure is provided with the *address* of the actual parameter (which must be a variable  $z$ ), and the formal parameter  $v$  is considered a synonym for  $z$ .<sup>2</sup>

Some subtleties of parameter passing come about as the result of *aliasing*. Consider the program on the right. If `<method>` is `val`, the program will print ‘4’. If it is `ref`, the program will print ‘8’. If both arguments were passed by value, the result would have been ‘3’.

```

proc main()
  int z;
  z := 3;
  p(z,z);
  write z;
end

proc p(ref int x, <method> int y)
  x := 4;
  y := y + x;
end

```

<sup>1</sup>Various instructions in the Brill assembly language can be used to take care of these rules for you.

<sup>2</sup>In terms of stack slots in the frame allocated for a procedure call, one slot is needed for a parameter passed by value. The parameter is simply treated as a local variable. A parameter passed by reference also needs one slot, but in this case, the *address* of the parameter is what is stored, and indirect addressing must be used to access the variable.

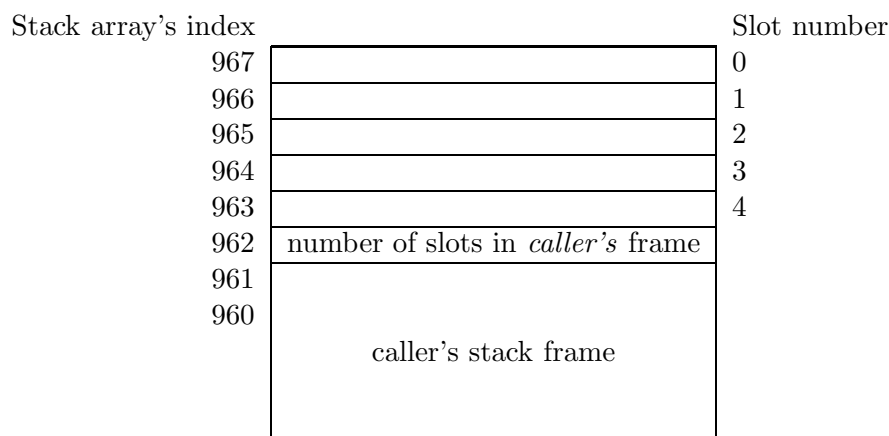


Figure 1: A stack frame on the ‘stack’ array

The rest of the semantics is the obvious—it follows standard conventions. For example, the execution of **while** *e* **do** *ss* **od** can be described as follows. First evaluate *e*. If *e* evaluates to **false**, the statement is equivalent to a no-op (a statement that does nothing). Otherwise the statement is equivalent to ‘*ss while e do ss od*’.

## The target language: Brill

Brill is an artificial target machine that closely resembles intermediate representations used by many compilers. An emulator for the Brill machine is supplied to you as a C program. This emulator reads Brill source files (which should be the output of your compiler) and executes them without further compilation. You are not required to modify the emulator, or understand its inner workings, and may treat it as a “black box” (although you may want to study the source code for your own benefit).

Brill has 1024 registers named *r0*, *r1*, *r2*, ... *r1023*. This is effectively an unlimited set of registers, and your compiler may treat it as such; your compiler may generate register numbers without checking whether they exceed 1023. Every register may contain a value of type **int** or **float** (referred to as **real** in the emulator).

Brill also has an area of main memory representing the stack, which contains zero or more stack frames. Each stack frame contains a number of stack slots. Each stack slot may contain a value of any type (it also holds type information about the value, for validation purposes), and you specify a stack slot by its stack slot number.

Figure 1 shows the “top” of the stack at some point. A stack pointer keeps track of the highest index used in the array, and stack slots can be accessed relative to this. Notice how the numbering of slots runs in the opposite direction of the array indices. In the example, the current stack frame (or activation record) has five slots (actually, it has one additional slot, used for remembering the size of the stack frame that will resume as current, once the active procedure returns).

An Brill program consists of a sequence of instructions, some of which may have labels. Although the emulator does not require it, good style dictates that each instruction should be on its own line. As in most assembly languages, you can attach a label to an instruction by preceding it with an identifier (the name of the label) and a colon. The label and the instruction

may be on the same line or different lines. Identifiers have the same format in Brill as in Snick.

The following lists the relevant opcodes of Brill (there are a few more), and for each opcode, shows how many operands it has, and what they are. *The destination operand is always the leftmost operand.*

push_stack_frame	framesize	
pop_stack_frame	framesize	
		# C analogues:
load	rI, slotnum	# rI = x
store	slotnum, rI	# x = rI
load_address	rI, slotnum	# rI = &x
load_indirect	rI, rJ	# rI = *rJ
store_indirect	rI, rJ	# *rI = rJ
int_const	rI, intconst	
real_const	rI, realconst	
string_const	rI, stringconst	
add_int	rI, rJ, rK	# rI = rJ + rK
add_real	rI, rJ, rK	# rI = rJ + rK
add_offset	rI, rJ, rK	# rI = rJ + rK
sub_int	rI, rJ, rK	# rI = rJ - rK
sub_real	rI, rJ, rK	# rI = rJ - rK
sub_offset	rI, rJ, rK	# rI = rJ - rK
mul_int	rI, rJ, rK	# rI = rJ * rK
mul_real	rI, rJ, rK	# rI = rJ * rK
div_int	rI, rJ, rK	# rI = rJ / rK
div_real	rI, rJ, rK	# rI = rJ / rK
cmp_eq_int	rI, rJ, rK	# rI = rJ == rK
cmp_ne_int	rI, rJ, rK	
cmp_gt_int	rI, rJ, rK	# etc.
cmp_ge_int	rI, rJ, rK	
cmp_lt_int	rI, rJ, rK	
cmp_le_int	rI, rJ, rK	
cmp_eq_real	rI, rJ, rK	
cmp_ne_real	rI, rJ, rK	
cmp_gt_real	rI, rJ, rK	
cmp_ge_real	rI, rJ, rK	
cmp_lt_real	rI, rJ, rK	
cmp_le_real	rI, rJ, rK	
and	rI, rJ, rK	# rI = rJ && rK
or	rI, rJ, rK	# rI = rJ    rK
not	rI, rJ	# rI = !rJ
int_to_real	rI, rJ	# rI = (float) rJ
move	rI, rJ	# rI = rJ

```

branch_on_true    rI, label          # if (rI) goto label
branch_on_false  rI, label          # if (!rI) goto label
branch_uncond    label              # goto label

call             label
call_builtin     builtin_function_name
return

debug_reg        rI
debug_slot       slotnum
debug_stack

halt

```

The `push_stack_frame` instruction creates a new stack frame. Its argument is an integer specifying how many slots the stack frame has; for example the instruction `stack_frame 5` creates a stack frame with five slots numbered 0 through 4. (In the emulator, it also reserves an extra slot, slot 5, to hold the size of the previous stack frame, for error detection purposes.)

The `pop_stack_frame` instruction deletes the current stack frame. Its argument is an integer specifying how many slots that stack frame has; it must match the argument of the `push_stack_frame` instruction that created the stack frame being popped.

The `load` instruction copies a value from the stack slot with the given number to the named register. The `store` instruction copies a value from the named register to the stack slot with the given number. The `load_address` instruction can be used by a caller to facilitate call by reference. The called procedure, having stored the address in the current stack frame, can then access and change the content of that address, by moving the address to a register and using `load_indirect` and `store_indirect`.

The `add_offset` and `sub_offset` instructions calculate addresses based on the offset from a given stack slot. They are useful when array components need to be accessed or updated. The instruction ‘`add_offset rI rJ rK`’ assumes that `rJ` holds an address, and `rK` holds an integer offset to be added to that address, the result being placed in `rI` (and similarly for `sub_offset`). Note that the Brill emulator is designed so that slot numbers grow in the opposite direction to how addresses grow, so `sub_offset` is appropriate when you want to *add* offsets to slot numbers, see Figure 1.

The `int_const`, `real_const`, and `string_const` instructions all load a constant of the specified type to the named register. The format of the constants is exactly the same as in Snick.

The `add_int`, `add_real`, `sub_int`, `sub_real`, `mul_int`, `mul_real`, `div_int` and `div_real` instructions perform arithmetic. The first part of the instruction name specifies the arithmetic operation, while the second part specifies the shared type of all the operands.

The `cmp_eq_int`, `cmp_ne_int`, `cmp_gt_int`, `cmp_ge_int`, `cmp_lt_int` and `cmp_le_int` instructions, and their equivalents for reals, perform comparisons, generating integer results. The middle part of the instruction name specifies the comparison operation, while the last part specifies the shared type of both input operands.

The `and`, `or` and `not` instructions each perform the “logical” operation of the same name.

The `int_to_real` instruction converts the integer in the source register to a real number in the destination register.



The `move` instruction copies the value in the source register (which may be of any type) to the destination register.

The `branch_on_true` instruction transfers control to the specified label if the named register contains a non-zero integer value. The `branch_on_false` instruction transfers control to the specified label if the named register contains 0. The `branch_uncond` instruction always transfers control to the specified label.

The `call` instruction calls the procedure whose code starts with the label whose name is the operand of the instruction, while the `call_builtin` instruction calls the built-in function whose name is the operand of the instruction. Procedures and functions take their first argument from register `r0`, their second from `r1`, and so on. During the call, the procedure may destroy the values in all the registers, so they contain nothing meaningful when the procedure returns. The exception is that the built-in functions that return a value, such as the read functions, put their return value in `r0` (see the example in Figure 3). When the called procedure executes the `return` instruction, execution continues with the instruction following the call instruction.

The following are all built-in functions: `read_int`, `read_real`, `read_bool`, `print_int`, `print_real`, `print_bool`, and `print_string`. (There are a few other built-ins that you will not need.) The read functions take no argument. They read a value of the indicated type (using `scanf`) from standard input, and return it in `r0`. The function `read_bool` accepts the strings “true” and “false”. The print functions take a single argument of the named type in `r0`, and print it to standard output; they return nothing.

Each `call` instruction pushes the return address (the address of the instruction following it) onto the stack. The `return` instruction transfers control to the address it pops off the stack.

The `halt` instruction stops the program.

Brill also supports comments, which start at a `#` character and continue until the end of the line. It may be useful to have the code generator insert comments, as in the example below.

The `debug_reg`, `debug_slot` and `debug_stack` instructions are Brill’s equivalent of debugging `printfs` in C programs: they print the value in the named register or stack slot or the entire stack. They are intended for debugging only; your submitted compiler should not generate them. If your code generator generates Brill code that does the wrong thing and you cannot sort out why, you can manually insert these instructions to better see what goes wrong. Calling the emulator with an `‘-i’` option gives a trace of execution.

Figure 2 shows the source program `gcd.snick`, and Figure 3 shows one possible translation. The Brill emulator starts execution with the first instruction in the program and stops when it executes the `halt` instruction. Note that the generated code therefore starts with a fixed two-instruction sequence that represents the Brill runtime system: the first instruction calls `main`, while the second (executed when `main` returns) is a `halt` instruction.

## Summary of Tasks, Suggestions

The compiler should take the name of a source file on the command line. It should write the corresponding target program to standard output, or report errors. The executable compiler must be called `snick`. On success, it must return 0 from `main`, and non-0 on failure.

You already have a working parser, and if not, you can use the supplied one (but in any case, correctness of the parser is your responsibility). There is no requirement to submit the pretty-printer, so it does not matter if you have to make changes to the AST that invalidate your pretty-printer.

```

proc main()
  int x;
  int y;
  int temp;
  int quotient;
  int remainder;

  write "Input two positive integers: ";

  read x;
  read y;

  write "\n";

  if x < y then
    temp := x;
    x := y;
    y := temp;
  fi

  write "The gcd of ";
  write x;
  write " and ";
  write y;
  write " is ";

  quotient := x / y;
  remainder := x - quotient * y;

  while remainder > 0 do
    x := y;
    y := remainder;
    quotient := x / y;
    remainder := x - quotient * y;
  od

  write y;
  write "\n";
end

```

Figure 2: The Snick program gcd.snick

```

        call proc_main
        halt
proc_main:
# prologue
    push_stack_frame 5
    int_const r0, 0
    store 0, r0    # int x
    store 1, r0    # int y
    store 2, r0    # int temp
    store 3, r0    # int quotient
    store 4, r0    # int remainder
# write
    string_const r0, "Input two positive integers: "
    call_builtin print_string
# read
    call_builtin read_int
    store 0, r0
# read
    call_builtin read_int
    store 1, r0
# write
    string_const r0, "\n"
    call_builtin print_string
# if
    load r0, 0
    load r1, 1
    cmp_lt_int r0, r0, r1
    branch_on_false r0, label0
# assignment
    load r0, 0
    store 2, r0
# assignment
    load r0, 1
    store 0, r0
# assignment
    load r0, 2
    store 1, r0
label0:
# write
    string_const r0, "The gcd of "
    call_builtin print_string
# write
    load r0, 0
    call_builtin print_int
# write
    string_const r0, " and "
    call_builtin print_string
# write
    load r0, 1
    call_builtin print_int

```

Figure 3: Translated program (first part)

```

# write
    string_const r0, " is "
    call_builtin print_string
# assignment
    load r0, 0
    load r1, 1
    div_int r0, r0, r1
    store 3, r0
# assignment
    load r0, 0
    load r1, 3
    load r2, 1
    mul_int r1, r1, r2
    sub_int r0, r0, r1
    store 4, r0
# while
label1:
    load r0, 4
    int_const r1, 0
    cmp_gt_int r0, r0, r1
    branch_on_false r0, label2
# assignment
    load r0, 1
    store 0, r0
# assignment
    load r0, 4
    store 1, r0
# assignment
    load r0, 0
    load r1, 1
    div_int r0, r0, r1
    store 3, r0
# assignment
    load r0, 0
    load r1, 3
    load r2, 1
    mul_int r1, r1, r2
    sub_int r0, r0, r1
    store 4, r0
    branch_uncond label1
label2:
# write
    load r0, 1
    call_builtin print_int
# write
    string_const r0, "\n"
    call_builtin print_string
# epilogue
    pop_stack_frame 5
    return

```

Figure 4: Translated program (remaining part)

The semantic analysis phase consists of a lot of checking that well-formedness conditions are met, as well as decorating of the AST with attributes that support code generation. The code generation phase consists of generating correct Brill code from the AST. Of these, well-formedness checking is arguably the part that has the lowest learning-outcome benefit for the time invested. The marking scheme encourages you to concentrate on code generation, and then deal with the correct handling of ill-formed programs as time allows.

You are encouraged to work stepwise and increase the part of the language covered as you go. It makes sense to write a module `symbol.ml` that offers the symbol table services. A module `analyze.ml` can do the semantic analysis of the AST, and it will want to store information in the symbol table. Work on getting the AST ready for code generation quickly—you can always add the well-formedness checks later, as time permits. A module `codegen.ml` can be responsible for code generation. It will also want to interact with the symbol table.

A possible approach to implementing `snick` incrementally is as follows (note that some static analysis is needed from the outset—types need to be determined for code generation):

1. Get the compiler working for the subset that consists of expressions (not including arrays) and the `write` statement. Assume that procedures do not take any arguments and cannot use recursion (no procedure calls), so that `main` works.
2. Add the `read` statement, and assignments.
3. Add compound statements (`if` and `while`). (Now you should be able to compile `gcd.snick`.)
4. Add procedure arguments and procedure calls, but initially for pass-by-value only.
5. Add reference parameters.
6. Add arrays.
7. Complete static analysis.

If you want to extend the task, there are two obvious improvements (and a bonus mark may apply for each):

- Add run-time checking for array bounds violations.
- Add optimisations. For example, a simple peephole analysis can be quite effective. The emulator will provide statistics if called with an ‘-s’ option.

The following directories are, or will be made, available on the LMS:

- `brill/` contains the Brill emulator. The make file will generate an executable called `brill`.
- `parser/` contains a `Snick` parser which is believed to be correct.
- `simple_tests/` contains a number small `snick` programs for testing.
- `contributed_tests/` will contain `snick` programs as submitted by teams (once they have been submitted).

## Procedure and assessment

The project may be solved in the teams, continuing from Stage 1. Each team should only submit once (under one of the members' name). If your team has changed since Stage 1, please let me know.

**By 5 May**, submit a single Snick program, which will be entered into a collection of test cases that will be made available to all. The program should be (syntactically, type, etc.) correct, but its runtime behaviour does not matter (whether it terminates, asks for input, divides by zero or whatever). Call your program `teamName.snick`, where `teamName` is your team's name, and submit a separate file `teamName.in` with the intended input to `teamName.snick`, if it requires input. For this stage, use `submit COMP90045 3a` to submit. It is possible to submit late, using `submit COMP90045 3a.late`, but late submissions will attract a penalty of 2 marks per calendar day late.

**By 19 May**, submit the code. There should be a `Makefile`, so that a `make` command generates `snick`. Do not submit the ML files that are generated by `ocamllex` and `ocamlyacc`; instead your `Makefile` should generate those ML files. Also, do not submit `brill.c` or other files related to Brill. For this last stage, use `submit COMP90045 3b` to submit.

This project counts for 14 of the 30 marks allocated to project work in this unit. Members of a group will receive the same mark, unless the group collectively sign a letter, specifying how the workload was distributed. We encourage the use of the LMS discussion forum and class time for discussions of ideas.

The marking sheet for Stage 3 will be made available on the LMS. Marks will be awarded on the basis of correctness (some 70%) and programming structure, style, readability, commenting and layout (some 30%). Out of the correctness marks, 60% will be directed towards code generation, with scanning, parsing, semantic analysis and symbol table handling counting for the remaining 40%.

## Appendix: Code format rules

Your OCaml programs should adhere with the following simple formatting rules:

- Each file should identify the team that produced it.
- Every non-trivial OCaml function should contain a comment at the beginning explaining its purpose and behaviour.
- Variable and function names must be meaningful.
- Significant blocks of code must be commented. However, not every statement in a program needs to be commented. Just as you can write too few comments, it is possible to write too many comments.
- Program blocks appearing in if-expressions, let clauses, etc., must be indented consistently. They can be indented using tabs or spaces, and can be indented 2, 4, or 8 spaces, as long as it is done consistently. Beware that some of the scaffolding code used spaces.
- Each program line should contain no more than 80 characters.

Graeme Gange  
25 April 2017